

Poincaré Embedding for Learning Hierarchical Representations

Marco Tsun Ting Lee

August 11, 2017

Overview

- 1 Background
 - Objective
 - Euclidean vs Hyperbolic
 - Paper's Result
- 2 Dataset
 - WordNet
- 3 Euclidean Space
 - Training
 - Evaluation
- 4 Hyperbolic Space
 - Definition & Properties
 - Training

What is a Word Embedding?

cats \rightarrow (0.8123, 0.7644, 82.2314, ...)

dogs \rightarrow (1.2323, 4.9712, 4.2790, ...)


king \rightarrow (6.1686, 78.1318, 6.4867, ...)

queen \rightarrow (34.1497, 6.4930, 3.2084, ...)


- Dimensionality Reduction (more efficient representation)
- Contextual Similarity (more expressive representation)
 - Words with similar context are clustered together.

Examples of Word Embedding


Traditional Example: $\Theta(\text{king}) - \Theta(\text{queen}) = \Theta(\text{man}) - \Theta(\text{woman})$



Male-Female

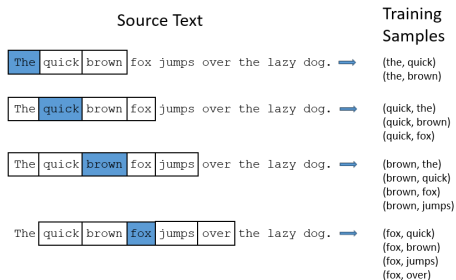


Verb tense



Country-Capital



Word2Vec - Existing Word Embedding Method



see <http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/>

Hierarchical Clustering - Method in Paper

Use Hierarchical Structure!



Hierarchical Clustering vs Word2Vec

Advantages:

- True context
- Rare words are not biased
- more efficient representation

Disadvantages:

- Most data do not have hierarchical structure
- Hard to get dataset with hierarchical structure built in

Why use Hyperbolic space?


Motivation:

Any finite subset of an hyperbolic space "looks like" a finite tree!

Informally, hyperbolic space can be thought of as a continuous version of tree.



Example

Imagine we have a tree...



Example

Assume each edge has length 1, we can define a discrete metric on the tree.



These two trees are the same!

Example

We only need the length between vertices to represent a tree!

```
# weighted graph adjacency matrix
#   A B C D E F G H I J K
A = [0, 1, 2, 2, 1, 2, 2, 3, 3, 1, 2]
B = [1, 0, 1, 1, 2, 3, 3, 4, 4, 2, 1]
C = [2, 1, 0, 2, 3, 4, 4, 5, 5, 3, 2]
D = [2, 1, 2, 0, 3, 4, 4, 5, 5, 3, 2]
E = [1, 2, 3, 3, 0, 1, 1, 2, 2, 2, 3]
F = [2, 3, 4, 4, 1, 0, 2, 3, 3, 3, 4]
G = [2, 3, 4, 4, 1, 2, 0, 1, 1, 3, 4]
H = [3, 4, 5, 5, 2, 3, 1, 0, 2, 4, 5]
I = [3, 4, 5, 5, 2, 3, 1, 2, 0, 4, 5]
J = [1, 2, 3, 3, 2, 3, 3, 4, 4, 0, 3]
K = [2, 1, 2, 2, 3, 4, 4, 5, 5, 3, 0]
graph = [A, B, C, D, E, F, G, H, I, J, K]
```




Euclidean vs Hyperbolic

Example

The optimal embedding in Euclidean space:

```
# weighted graph adjacency matrix
# A B C D E F G H I J K
A = [0, 1, 2, 2, 1, 2, 2, 3, 3, 1, 2]
B = [1, 0, 1, 1, 2, 3, 3, 4, 4, 2, 1]
C = [2, 1, 0, 2, 3, 4, 4, 5, 5, 3, 2]
D = [2, 1, 2, 0, 3, 4, 4, 5, 5, 3, 2]
E = [1, 2, 3, 3, 0, 1, 1, 2, 2, 2, 3]
F = [2, 3, 4, 4, 1, 0, 2, 3, 3, 3, 4]
G = [2, 3, 4, 4, 1, 2, 0, 1, 1, 3, 4]
H = [3, 4, 5, 5, 2, 3, 1, 0, 2, 4, 5]
I = [3, 4, 5, 5, 2, 3, 1, 2, 0, 4, 5]
J = [1, 2, 3, 3, 2, 3, 3, 4, 4, 0, 3]
K = [2, 1, 2, 2, 3, 4, 4, 5, 5, 3, 0]
graph = [A, B, C, D, E, F, G, H, I, J, K]
```



```
# euclid dist adjacency matrix
# A B C D E F G H I J K
A = [0, 1.1, 2.1, 2.3, 1.1, 2, 2.1, 3.2, 3.2, 1.4, 2.2]
B = [1.1, 0, 1.3, 1.3, 2.1, 3.1, 3, 3.9, 4.2, 2, 1.1]
C = [2.1, 1.3, 0, 2, 3.1, 3.8, 4.2, 5.2, 5.2, 2.1, 1.1]
D = [2.3, 1.3, 2, 0, 3.1, 4.3, 3.7, 4.3, 4.9, 3.3, 1]
E = [1.1, 2.1, 3.1, 3.1, 0, 1.4, 1.1, 2.3, 2.1, 2, 3.2]
F = [2, 3.1, 3.8, 4.3, 1.4, 0, 1.9, 3, 2.2, 1.9, 4.2]
G = [2.1, 3, 4.2, 3.7, 1.1, 1.9, 0, 1.2, 1.2, 3.1, 4.1]
H = [3.2, 3.9, 5.2, 4.3, 2.3, 3, 1.2, 0, 1.3, 4.3, 4.8]
I = [3.2, 4.2, 5.2, 4.9, 2.1, 2.2, 1.2, 1.3, 0, 3.9, 5.2]
J = [1.4, 2, 2.1, 3.3, 2, 1.9, 3.1, 4.3, 3.9, 0, 2.8]
K = [2.2, 1.1, 1.1, 1, 3.2, 4.2, 4.1, 4.8, 5.2, 2.8, 0]
```


L2 Loss = 7.166

see <https://github.com/marcoleewow/Find-Optimal-Space-Embedding-for-Trees>

Example

The optimal embedding in Hyperbolic space:

```
# weighted graph adjacency matrix
# A B C D E F G H I J K
A = [0, 1, 2, 2, 1, 2, 2, 3, 3, 1, 2]
B = [1, 0, 1, 1, 2, 3, 3, 4, 4, 2, 1]
C = [2, 1, 0, 2, 3, 4, 4, 5, 5, 3, 2]
D = [2, 1, 2, 0, 3, 4, 4, 5, 5, 3, 2]
E = [1, 2, 3, 3, 0, 1, 1, 2, 2, 2, 3]
F = [2, 3, 4, 4, 1, 0, 2, 3, 3, 3, 4]
G = [2, 3, 4, 4, 1, 2, 0, 1, 1, 3, 4]
H = [3, 4, 5, 5, 2, 3, 1, 0, 2, 4, 5]
I = [3, 4, 5, 5, 2, 3, 1, 2, 0, 4, 5]
J = [1, 2, 3, 3, 2, 3, 3, 4, 4, 0, 3]
K = [2, 1, 2, 2, 3, 4, 4, 5, 5, 3, 0]
graph = [A, B, C, D, E, F, G, H, I, J, K]
```



```
# hyp dist adjacency matrix
# A B C D E F G H I J K
A = [0. 1.1 2.1 2.1 1.1 2. 2.1 3. 3.1 1.3 2. ]
B = [1.1 0. 1.1 1.3 2.1 3.1 3. 3.9 4.1 2. 1.3]
C = [2.1 1.1 0. 1.5 3.1 4.2 4.1 5. 5.1 3. 1.5]
D = [2.1 1.3 1.5 0. 3. 4. 3.8 4.7 4.9 3.2 2.4]
E = [1.1 2.1 3.1 3. 0. 1.2 1.1 2.1 2.1 2. 3.1]
F = [2. 3.1 4.2 4. 1.2 0. 2. 3.1 2.7 2.6 4. ]
G = [2.1 3. 4.1 3.8 1.1 2. 0. 1.2 1.1 3.1 4. ]
H = [3. 3.9 5. 4.7 2.1 3.1 1.2 0. 1.8 4.1 5. ]
I = [3.1 4.1 5.1 4.9 2.1 2.7 1.1 1.8 0. 4. 5.1]
J = [1.3 2. 3. 3.2 2. 2.6 3.1 4.1 4. 0. 2.6]
K = [2. 1.3 1.5 2.4 3.1 4. 4. 5. 5.1 2.6 0. ]
```

L2 Loss = 1.928!

see <https://github.com/marcoleewow/Find-Optimal-Space-Embedding-for-Trees>

Quasi-Isometry

In math terms,

Theorem

$\forall n \in \mathbb{Z}, \delta > 0, \exists$ a constant C s.t. the following holds:

if x_1, \dots, x_n are points in a δ -hyperbolic space X ,

\exists a finite tree T and an embedding $f : T \rightarrow X$ s.t.

$x_i \in f(T) \forall i = 1, \dots, n$ and

$$\forall i, j : d_T(f^{-1}(x_i), f^{-1}(x_j)) \leq d_X(x_i, x_j) \leq d_T(f^{-1}(x_i), f^{-1}(x_j)) + C$$

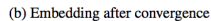
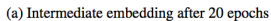
The constant C can be taken to be $\delta \cdot h(n)$ with $h(n) = O(\log n)$ and this is optimal!

Hyperbolic vs Euclidean Results

		Dimensionality					
		5	10	20	50	100	200
Euclidean	Rank	3311.1	2199.5	952.3	351.4	190.7	81.5
	MAP	0.024	0.059	0.176	0.286	0.428	0.490
Poincaré	Rank	5.7	4.3	4.9	4.6	4.6	4.6
	MAP	0.825	0.852	0.861	0.863	0.856	0.855

Hyperbolic Advantages

- Outperform Euclidean embedding significantly
- Much more efficient representation!
- Did not add much computational intensity



WordNet

- Clear latent hierarchical structure
- Easy to use
- Publicly available for download
- Enough data for experiments

```

>>> wn.synset('dog.n.01')
Synset('dog.n.01')
>>> print(wn.synset('dog.n.01').definition())
a member of the genus Canis (probably descended from the common wolf) that has been domesticated by man since prehistoric times; occurs in many breeds
>>> len(wn.synset('dog.n.01').examples())
1
>>> print(wn.synset('dog.n.01').examples()[0])
the dog barked all night
>>> wn.synset('dog.n.01').lemmas()
[Lemma('dog.n.01.dog'), Lemma('dog.n.01.domestic_dog'), Lemma('dog.n.01.Canis_familiaris')]
>>> [str(lemma.name()) for lemma in wn.synset('dog.n.01').lemmas()]
['dog', 'domestic_dog', 'Canis_familiaris']
>>> wn.lemma('dog.n.01.dog').synset()
Synset('dog.n.01')

```

<http://www.nltk.org/howto/wordnet.html>

All Noun Synsets

```
>>> for synset in list(wn.all_synsets('n'))[:10]:  
...     print(synset)  
...  
Synset('entity.n.01')  
Synset('physical_entity.n.01')  
Synset('abstraction.n.06')  
Synset('thing.n.12')  
Synset('object.n.01')  
Synset('whole.n.02')  
Synset('congener.n.03')  
Synset('living_thing.n.01')  
Synset('organism.n.01')  
Synset('benthos.n.02')
```

Hypernymy Relation

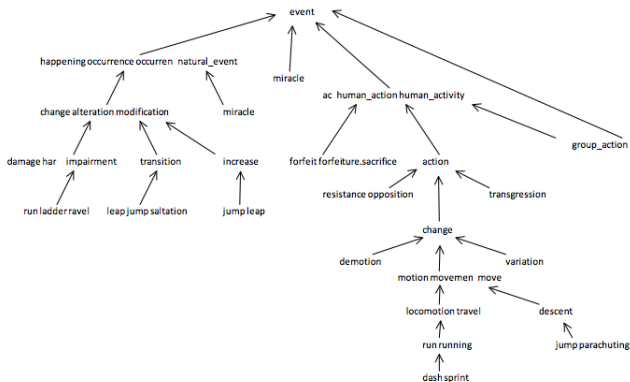
Hypernymy relation = 'is-a' relation

e.g. dog (hyponym) 'is-a' canine (hypernym)

Compute transitive closures of synsets

```
>>> dog = wn.synset('dog.n.01')
>>> hypo = lambda s: s.hyponyms()
>>> hyper = lambda s: s.hypernyms()
>>> list(dog.closure(hypo, depth=1)) == dog.hyponyms()
True
>>> list(dog.closure(hyper, depth=1)) == dog.hypernyms()
True
>>> list(dog.closure(hypo))
[Synset('basenji.n.01'), Synset('corgi.n.01'), Synset('cur.n.01'),
 Synset('dalmatian.n.02'), Synset('great_pyrenees.n.01'),
 Synset('griffon.n.02'), Synset('hunting_dog.n.01'), Synset('lapdog.n.01'),
 Synset('leonberg.n.01'), Synset('mexican_hairless.n.01'),
 Synset('newfoundland.n.01'), Synset('pooch.n.01'), Synset('poodle.n.01'), ...]
>>> list(dog.closure(hyper))
[Synset('canine.n.02'), Synset('domestic_animal.n.01'), Synset('carnivore.n.01'),
 Synset('animal.n.01'), Synset('placental.n.01'), Synset('organism.n.01'),
 Synset('mammal.n.01'), Synset('living_thing.n.01'), Synset('vertebrate.n.01'),
 Synset('whole.n.02'), Synset('chordate.n.01'), Synset('object.n.01'),
 Synset('physical_entity.n.01'), Synset('entity.n.01')]
```

Directed Acyclic Graph



WordNet Data Samples

- Total of 82,115 nouns and 743,241 hypernymy relations.
- Index the nouns: $\{0:\text{cat}, 1:\text{dog}, 2:\text{canine}, \dots, 82114:\text{queen}\}$
- Let $\mathcal{D} = \{(u, v)\}$ be the set of observed hypernymy relations between noun pairs.
- e.g. dog 'is-a' canine $\Rightarrow (1, 2) \in \mathcal{D}$

Objective

- Euclidean Distance $d(\mathbf{u}, \mathbf{v}) = \|\mathbf{u} - \mathbf{v}\|^2$
- Unit ball $\mathcal{B}^d = \{x \in \mathbb{R}^d : \|x\| < 1\}$

Let n be total number of nouns (in our case $n = 82115$).

We want to find $\Theta = \{\theta_i\}_{i=1}^n$, where $\theta_i \in \mathcal{B}^d$

To estimate Θ , we solve the optimization problem:

$$\Theta' \leftarrow \operatorname{argmin} \mathcal{L}(\Theta)$$

Negative Sampling Loss


$$\mathcal{L}(\Theta) = \sum_{(u,v) \in \mathcal{D}} \log \frac{e^{-d(u,v)}}{\sum_{v' \in \mathcal{N}(u)} e^{-d(u,v')}}}$$

, where $\mathcal{N}(u) = \{v | (u, v) \notin \mathcal{D}\}$


- soft ranking loss
- stochastically sample 10 negative samples
- differentiable
- update u only! All v 's are fixed

Loss Function Visualization


Initialization: $\Theta_{ij} \sim \mathcal{U}(-0.001, 0.001)$



Loss Function Visualization



Loss Function Visualization



SGD Optimizer

$$\text{proj}(\theta) = \begin{cases} \frac{\theta}{\|\theta\|} - \epsilon & \text{if } \|\theta\| \geq 1 \\ \theta & \text{otherwise.} \end{cases}$$

$$\theta_{t+1} \leftarrow \text{proj}(\theta_t - \eta_t \nabla_E \mathcal{L})$$

Link Prediction

- split data samples into 80% train, 10% valid and 10% test samples.
- Test generalization performance by hiding some links and check if we can predict them using distance in embedded space.

Rank Metric

For each pair of $(u, v) \in \mathcal{D}$, We rank its distance $d(\mathbf{u}, \mathbf{v})$ among the ground truth negative samples.

Mean Average Precision (MAP)

Rank nodes by closest distance to u , and see if they are linked.

e.g.

number of links of $u = 4$.

links: 1, 0, 1, 1, 0, 1.

Precision: $1/1, 1/2, 2/3, 3/4, 3/5, 4/6$

Average precision (take average across corrected links):

$(1/1 + 2/3 + 3/4 + 4/6)/4 = 0.77$

MAP: Take mean across all samples.


Results

		Dimensionality					
		5	10	20	50	100	200
Euclidean	Rank	3311.1	2199.5	952.3	351.4	190.7	81.5
	MAP	0.024	0.059	0.176	0.286	0.428	0.490

Why Not Euclidean?

- Curse of dimensionality
- Inefficient representation

Riemannian Manifold Examples



Poincaré Ball Model \mathbb{B}^d

Poincaré Ball Model \mathbb{B}^d :


$$\mathbb{B}^d = \{\mathbf{x} \in \mathbb{R}^d : \|\mathbf{x}\| < 1\}$$

equipped with hyperbolic metric:

$$d(\mathbf{u}, \mathbf{v}) = \operatorname{arcosh}\left(1 + 2 \frac{\|\mathbf{u} - \mathbf{v}\|^2}{(1 - \|\mathbf{u}\|^2)(1 - \|\mathbf{v}\|^2)}\right),$$


where $\mathbf{u}, \mathbf{v} \in \mathbb{B}^d$.

Poincaré Ball Model \mathbb{B}^2




Hyperbolic Space Properties

Parallel Postulate




Hyperbolic Space Properties

Parallel Postulate



Hyperbolic Space Properties

Parallel Postulate



Hyperbolic Space Properties

Circle of radius r :

$$\text{Circumference} = 2\pi \sinh r > 2\pi r \text{ for } r > 0.$$

$$\text{Area} = 2\pi \cosh r - 1 > \pi r^2 \text{ for } r > 0.$$

Ball of radius r :


$$\text{Volume} = \pi(\sinh(2r) - 2r) > \frac{4}{3}\pi r^3 \text{ for } r > 0.$$

Intuition: "Stores more information" locally

Riemannian Gradient

$$\nabla_g = g_\theta^{-1} \nabla_E,$$

$$\text{where } g_\theta^{-1} = \frac{(1 - \|\theta\|^2)^2}{4}$$



ignore notations in the above diagram!

Riemannian SGD

Full update equation:

$$\theta_{t+1} \leftarrow \text{proj}(\theta_t - \eta_t \nabla_g \mathcal{L})$$

$$= \theta_{t+1} \leftarrow \text{proj}(\theta_t - \eta_t g_{\theta_t}^{-1} \nabla_E \mathcal{L})$$

with hyperbolic distance for \mathcal{L} instead of Euclidean!

Results

		Dimensionality					
		5	10	20	50	100	200
Euclidean	Rank	3311.1	2199.5	952.3	351.4	190.7	81.5
	MAP	0.024	0.059	0.176	0.286	0.428	0.490
Poincaré	Rank	5.7	4.3	4.9	4.6	4.6	4.6
	MAP	0.825	0.852	0.861	0.863	0.856	0.855

The End

